A Simple Divide-and-Conquer Algorithm for Constructing Delaunay Triangulations in O(n log log n) Expected Time

Rex A. Dwyer[*]

Computer Science Department
Carnegie-Mellon University
Pittsburgh, Pennsylvania 15213

0. Abstract.

We present a modification to the divide-and-conquer algorithm of Guibas & Stolfi [GS] for computing the Delaunay triangulation of n sites in the plane. The change reduces its $\Theta(n \log n)$ expected running time to $O(n \log \log n)$ for a large class of distributions which includes the uniform distribution in the unit square. The modified algorithm is significantly easier to implement than the optimal linear-expected-time algorithm of Bentley, Weide & Yao [BWY]. Unlike the incremental methods of Ohya, Iri & Murota [OIM] and Maus [M] it has optimal $O(n \log n)$ worst-case performance. The improvement extends to the computation of the Delaunay triangulation in the L_p metric for $1 < p \leq \infty$. Experimental evidence presented demonstrates that in the Euclidean case the modified algorithm performs very well for $n \leq 2^{16}$, the range of the experiments. We conjecture that its average running time is no more than twice optimal for n less than seven trillion.

1. Introduction.

The *Voronoi diagram* generated by a set $S = \{P_1, P_2, \ldots, P_n\}$ of n points in the plane (called *sites*) partitions the plane into n convex regions $\mathcal{V}_i = \{ P \mid \forall j : d(P, P_i) \leq d(P, P_j) \}$. Each region contains the points lying nearer to the site in its interior than to any other site.

The straight-line dual of the Voronoi diagram is called the *Delaunay triangulation*, denoted $DT(S)$. P_i and P_j are adjacent in $DT(S)$ if and only if $\mathcal{V}_i$ and $\mathcal{V}_j$ share a common edge. *(Fig. 1.1.)* If, as is assumed in the sequel, no four sites are cocircular, the Delaunay triangulation partitions the convex hull of S into triangles. The Voronoi diagram can be constructed from the Delaunay triangulation in $O(n)$ time and *vice versa*.

Existing algorithms for constructing the widely-applied Voronoi diagram and Delaunay triangulation fall into two categories: divide-and-conquer algorithms and incremental algorithms.

Shamos [Sh] presents the first divide-and-conquer algorithm and shows its $O(n \log n)$ worst-case running time to be optimal under the real RAM model of computation. Lee & Schachter [LS] describe a dual algorithm for constructing the Delaunay triangulation. Guibas & Stolfi [GS] advocate the Delaunay triangulation as an intermediate step in the construction of the Voronoi diagram, and present ideal data structures for the problem. Hwang [H], Lee & Wang [LW], and Lee [L] present $O(n \log n)$ divide-and-conquer algorithms for the Voronoi diagram in the L_1, L_1 and L_∞, and general L_p metrics respectively. Recent work of Ohya, Iri & Murota [OIM] shows the average running time of the Euclidean-metric algorithms to be $\Omega(n \log n)$ when the sites are uniformly distributed in the unit square. Their proof extends naturally to the L_p metrics.

Various incremental algorithms, which con-

Fig. 1.1. A Voronoi diagram and its dual.

struct the Voronoi diagram by adding new sites one by one, differ principally in the order in which sites are added. Sibson & Green [SG] were early advocates of such a method, giving an algorithm with $O(n^2)$ worst-case and $O(n^{3/2})$ average running times. The algorithms of Ohya, Iri & Murota [OIM] and Maus [M] achieve optimal linear average time but only quadratic worst-case time. Fortune's new sweepline algorithm [F] uses a clever geometric transformation to achieve $O(n \log n)$ worst-case time; only the trivial bounds on its average performance are known.

Bentley, Weide & Yao [BWY] propose a complicated hybrid algorithm which achieves optimal running time in both the worst-case and expected senses. It invokes a divide-and-conquer algorithm on the "outer" sites lying near the boundary of the unit square. The Voronoi polygon of each "inner" site is constructed in constant expected time by a "spiral search" among its neighbors. In the unlikely event that more than $O(\log n)$ time is expended on any inner site, spiral search is abandoned and the divide-and-conquer algorithm is applied to the entire set of sites.

We present an easily implemented improvement to the divide-and-conquer algorithm of Guibas & Stolfi. If the floor function can be computed in constant time, the change maintains its $O(n \log n)$ worst-case complexity, but lowers its average-case complexity to $O(n \log \log n)$ when the sites are drawn independently from any of a large class of distributions which includes the uniform distribution on the unit square. The modification can be made to any of the divide-and-conquer algorithms including those for the

L_p metric for $1 < p \leq \infty$. Our experimental evidence demonstrates that in the Euclidean metric the improved algorithm performs very well for $n \leq 2^{16}$, the range of the experiments. We conjecture that its average running time is no more than twice optimal for n less than seven trillion.

2. Preliminaries.

We now restate some useful facts about Delaunay triangulations and Voronoi diagrams and review the results of Guibas & Stolfi upon which our work depends.

Lemma 2.1. *Every triangulation of a set of n sites of which k lie on the convex hull (i.e., on the boundary of the convex hull) has $3(n-1) - k$ edges and $2(n-1) - k$ triangles.*

Proof. Let t be the number of triangles and e the number of edges. All the edges belong to two triangles except for the k on the convex hull, thus $2e - k = 3t$. The results are obtained by solving this equation for t or e and substituting into $n - e + (t + 1) = 2$ (Euler's formula). $\square$

Corollary 2.2. *At most $3n$ non-intersecting edges can be constructed on a set of n sites.*
Proof. Any set of non-intersecting edges can be extended to form a triangulation. $\square$

Lemma 2.3 *If sites P_i, P_j and P_k are vertices of a triangle in the Delaunay triangulation, the circle passing through these three sites contains no other sites.*

Proof. The center of the circle is equidistant from the three sites and is on the boundary of the Voronoi region of each of them. Thus it cannot lie nearer to a fourth site than to P_i, P_j and P_k. $\square$

Corollary 2.4. *If sites P_i and P_j are endpoints of a Delaunay edge, there is some circle containing no sites which passes through P_i and P_j.*

Proof. Each edge is a side of some triangle. $\square$

Corollary 2.5. *Every convex hull edge is a Delaunay edge.*

Proof. One of the half-planes defined by the affine extension of the edge is a degenerate site-free circle. $\square$

Sibson [S] has showed that, degenerate cases with four cocircular sites excepted, the Delaunay

277

triangulation is the only triangulation satisfying Lemma 2.3.

Guibas & Stolfi's algorithm first constructs the Delaunay triangulation. The Voronoi diagram is then found in linear time. The Delaunay triangulation is constructed as follows:

1) The sites are sorted along the x-axis.
2) If there are three or fewer sites, the Delaunay triangulation is constructed directly. Otherwise, the sites are divided into two approximately equal sets by a line perpendicular to the x-axis, Step 2 is recursively applied to construct the Delaunay triangulations of these sets, and the results are merged.

The merge procedure forms the foundation of our own algorithm as well as of Guibas & Stolfi's. Let ℓ be the dividing line of Step 2, and let $\mathcal{L}$ and $\mathcal{R}$ be the sets lying to the left and the right of ℓ. Clearly two edges of the convex hull of $\mathcal{L} \cup \mathcal{R}$ cross ℓ. Merging begins with a search for the endpoints of the lower of the two. The search begins with the sites of $\mathcal{L}$ and $\mathcal{R}$ lying nearest ℓ and alternately advances clockwise around the convex hull of $\mathcal{L}$ and counterclockwise around the convex hull of $\mathcal{R}$. Once its endpoints are found and the lower hull edge is created, the other new edges are created in the order in which they cross ℓ. Old edges are removed as necessary. The essential features of the merge procedure are summarized in the following theorem:

Theorem 2.6. Let ℓ, $\mathcal{L}$ and $\mathcal{R}$ be as above. When merging $DT(\mathcal{L})$ and $DT(\mathcal{R})$ to construct $DT(\mathcal{L} \cup \mathcal{R})$:

a) Only edges joining two sites in $\mathcal{L}$ and edges joining two sites in $\mathcal{R}$ are deleted.
b) Only edges crossing ℓ and joining a site in $\mathcal{L}$ to one in $\mathcal{R}$ are created.
c) The worst-case running time of the merge is bounded by a function linear in the sum of three components:
　i) the number of sites examined to find the endpoints of the lower of the two edges of the convex hull of $\mathcal{L} \cup \mathcal{R}$ which cross ℓ.
　ii) the number of edges deleted, and
　iii) the number of edges created.
d) The worst-case running time of the merge is $O(|\mathcal{L} \cup \mathcal{R}|)$.

Proof. Guibas & Stolfi. □

Our analysis requires a more refined bound on the running time of the merge than those of Theorem 2.6.

Theorem 2.7. *The total running time of an algorithm based on Guibas & Stolfi's merge procedure is bounded by a linear function of the number of edges created.*

Proof. By Theorem 2.6.c, it suffices to show that the number of edges deleted and the number of sites examined to construct the lower convex hull edge are bounded by the number of edges created in the merge.

Clearly no edge can be deleted which has not first been created during a previous invocation of the merge procedure.

It is also true that each of the sites examined in the convex-hull-edge construction of a particular invocation of the merge procedure receives a new edge during the *same* invocation. Since each one falls on the boundary of the convex hull of $\mathcal{L}$ or $\mathcal{R}$, it must be the vertex of some angle larger than 180 degrees before the merge. However, it must also lie in the interior of the convex hull of $\mathcal{L} \cup \mathcal{R}$, and it therefore cannot be the vertex of such a large angle after the merge. □

3. The Modified Algorithm and Analysis of Its Worst-Case Running Time.

The unit square is partitioned into $n/\log n$ square cells with sides of length $\sqrt{\log n / n}$. The Delaunay triangulation of the sites within each cell is constructed with the Guibas-Stolfi algorithm. The triangulations within each row of cells are merged in pairs until the triangulation of the row has been completed. Then row triangulations are merged in pairs to complete the triangulation of the entire set of sites. *(Fig. 3.1.)*

Theorem 3.1 Algorithm A uses $O(n \log n)$ time in the worst case.

Proof. Step 0 requires $O(n)$ time, since the floor function is assumed to be computable in constant time.

Number the cells arbitrarily and let n_i be the number of sites in the i-th cell. Since Guibas & Stolfi's algorithm requires $O(n \log n)$ time in the worst case, Step 1 can be completed in time $\sum_i O(n_i \log n_i) \le \sum_i O(n_i \log n) \le O(n \log n)$.

278

```
{ Step 0: Sort Sites into Buckets }
m := √n/log n
for P ∈ S do insert P into B⌊mxₚ⌋,⌊myₚ⌋
{ Step 1: Triangulate Cells }
for i := 0 to m − 1 do
    for j := 0 to m − 1 do
        DTᵢⱼ := Guibas_Stolfi_DT(Bᵢⱼ)
{ Step 2: Merge Cells into Rows }
for k := 0 to lg m do
    for i := 0 to m − 1 do
        for j := 0 to m − 1 by 2^(k+1) do
            DTᵢⱼ := merge(DTᵢⱼ, DT_(i,j+2ᵏ));
{ Step 3: Merge Rows }
for k := 0 to lg m do
    for i := 0 to m − 1 by 2^(k+1) do
        DTᵢ₀ := merge(DTᵢ₀, DT_(i+2ᵏ,0));
return DT₀₀;
```

Figure 3.1. Algorithm A

In Step 2, no site is involved in more than a single merge for any fixed value of k. Since by Corollary 2.2, the number of points involved in the merges bounds the number of edges created, and by Theorem 2.7 the number of edges created bounds the running time, no more than $O(n)$ time is required for each iteration of the k loop. The k loop is executed $\lg(\sqrt{n/\log n}) < \lg n$ times, thus $O(n \log n)$ time is required for Step 2.

Step 3 can be handled exactly like Step 2. Summing the times for the three steps, we find that at most $O(n \log n)$ time is required. □

4. Analysis of Expected Time.

We will call a distribution with density function f *quasi-uniform* in a region if, for some strictly positive constants c_1 and c_2, $f(x,y) = 0$ outside the region and $c_1 \leq f(x,y) \leq c_2$ inside the region. By modifying constants in the proofs which follow, it is not difficult to show that the expected running time of Algorithm A is $O(n \log \log n)$ for any quasi-uniform distribution in a rectangle. However, we restrict our attention to the uniform distribution in the unit square for the sake of expository simplicity.

During Step 1, we accept the $O(n \log n)$ worst case performance of Guibas & Stolfi's algorithm to construct Delaunay triangulations within each cell. We need only take care to show that it is

unlikely that very many sites fall within a single cell, making worst-case performance for that cell unacceptably large.

We next show that when merging two Delaunay triangulations most sites to which new edges are constructed must lie near the boundaries of the two triangulations. New edges to sites far from the boundaries must necessarily be so long that they determine large circles likely to contain other sites, in violation of Corollary 2.4. Since the new edges are non-intersecting, by Corollary 2.2 and Theorem 2.7 the number of sites to which new edges are constructed bounds the running time.

In the sequel we assume that $\mathcal{U}$ is the unit square and that $\mathcal{S} = \{P_1, P_2, \ldots, P_n\}$ is a set of n sites chosen independently from the uniform distribution on $\mathcal{U}$.

Lemma 4.1. The probability that any fixed cell contains more than $e \log n$ sites is less than $1/n$, where e is the base of the natural logarithms.

Proof. Since the cells are squares with sides of length $\sqrt{\log n/n}$, the probability that a given site lies in the cell is equal to $\log n/n$, the area of the cell. Call this probability p, and let N be the number of sites in the cell. Then N has a binomial distribution, and

$$\Pr\{N \geq e \log n\} = \sum_{k \geq e \log n} \Pr\{N = k\}$$

$$\leq \sum_{k \geq e \log n} e^{(k - e \log n)} \Pr\{N = k\}$$

$$\leq n^{-e} \sum_{k \geq 0} e^k \Pr\{N = k\}$$

$$= n^{-e} \sum_{k \geq 0} \binom{n}{k} (ep)^k (1-p)^{n-k}$$

$$= n^{-e}(ep + 1 - p)^n$$

$$= n^{-e} \exp(n \log(1 + p(e - 1))).$$

Since $\log(1 + x) < x$ for $x > 0$, it follows that $\Pr\{N \geq e \log n\} < n^{-e} \exp(np(e - 1))$ $= n^{-e} \exp((\log n)(e - 1)) = n^{-e} \cdot n^{e-1} = 1/n$. □

In passing, we note that the same argument can be used to show that the probability that a fixed cell contains more than $(e + m - 1) \log n$ sites is at most $1/n^m$.

Theorem 4.2. *Step 1 of Algorithm A requires $O(n \log \log n)$ expected time.*

Proof. Since there are $n/\log n$ cells, Lemma 4.1 implies that $\Pr\{\text{some cell has} \geq e \log n \text{ sites}\} \leq (n/\log n) \cdot \Pr\{\text{a fixed cell has} \geq e \log n \text{ sites}\} \leq 1/\log n$. Since even these inputs can be handled in $O(n \log n)$ time, their contribution to the expected running time of Step 1 is at most $O\left((n \log n) \cdot (1/\log n)\right) = O(n)$. For the other cases, at most $n/\log n$ subproblems of size at most $e \log n$ must be solved using the $O(n \log n)$ worst-case algorithm of Guibas & Stolfi, making the total expected time $O\left((n/\log n)e \log n \log(e \log n)\right)$ or $O(n \log \log n)$ for Step 1. $\square$

We now bound the expected time of the merging steps. We first show that most new edges lie in a region near the boundaries of the diagrams being merged. We then bound the number of edges which we expect to create in these boundary regions in all the merges of Steps 2 and 3.

The following lemma will be proved in a more general setting in the next section:

Lemma 4.3. *Let P, Q, and P' be three points lying along a line in that order. Let M and N be the other vertices of the square with diagonal PQ. Then any circle passing through P and P' completely encloses either $\triangle PQM$ or $\triangle PQN$ or both.*

Lemma 4.4. *Let $p = 2 \log \log n/n$, let T be a subset of $\mathcal{U}$ with area exceeding p, and let S' be a set of at least $n - 2$ sites drawn independently from the uniform distribution on $\mathcal{U}$. Then $\Pr\{S' \cap T = \emptyset\} = O(1/\log n)$.*

Proof.

$$\Pr\{S' \cap T = \emptyset\} \leq (1 - p)^{n-2}$$
$$= \exp\left((n - 2)\log(1 - p)\right)$$
$$< \exp\left(-p(n - 2)\right)$$
$$\leq \exp\left((-\log \log n)(4/n - 2)\right)$$
$$= \log n^{(4/n)-2} = O(1/\log n). \; \square$$

Theorem 4.5. *Let $\square ABCD$ and $\square CDEF$ lie inside $\mathcal{U}$, let $\mathcal{R} = S \cap \square ABCD$ and $\mathcal{L} = \widehat{} \cap \square CDEF$, and let $h = |CD|$. Then there is a merge region $M \subset \square ABCD$ such that:*

a) $\text{Area}(M) = O(h\sqrt{\log \log n/n} + \log n \log \log n/n)$

b) if $P \in \mathcal{R}$ but $P \notin M$ then $p = \Pr\{\exists P' \in \mathcal{L} \mid (P, P') \in DT(\mathcal{L} \cup \mathcal{R})\} = O(1/\log n)$.

Proof. For convenience, we use a coordinate system with origin at C, B on the positive x-axis, and D on the positive y-axis. *(Fig. 4.1.)* We now verify that conditions a) and b) are satisfied by

$$M = \square ABCD \cap \{(x,y) \mid (x \leq 8\sqrt{\log \log n/n})$$
$$\lor(xy \leq 24 \log \log n/n)$$
$$\lor(x(h - y) \leq 24 \log \log n/n)\}$$

Fig. 4.1. The merge region M.

Since the two hyperbolas defining M are mirror images, and since $|BC| \leq 1$, we have

$$\text{Area}(M) \leq 8h\sqrt{\log \log n/n}$$
$$+ 2\int_{8\sqrt{\log \log n/n}}^{1} \frac{24 \log \log n}{nx} dx$$
$$= O(h\sqrt{\log \log n/n} + \log n \log \log n/n)$$

Now let $P = (x_P, y_P)$ be a site in $\mathcal{R}$ but outside M, let P' be a site in $\mathcal{L}$, and let Q be the point of intersection of PP' and CD. Let $\square PMQN$ be the square with M above and N below PQ. Further suppose that P and P' are joined in $DT(\mathcal{L} \cup \mathcal{R})$. By Corollary 2.4, there is a circle C through P and P' whose intersection with $\square ABCD$ is site-free. Then according to Lemma 4.3 either $(\triangle PQM \cap \square ABCD)$ or $(\triangle PQN \cap \square ABCD)$ is site-free. We will show that each of $\triangle PQM$ and $\triangle PQN$ completely contains one of ten regions in $\square ABCD$, each of which has area exceeding $2 \log \log n/n$. The conditions of Lemma 4.4 are satisfied for each region with $S' = S - \{P, P'\}$, so the probability that any fixed region is site-free is $O(1/\log n)$. Therefore the probability

that at least one of the ten regions is site-free is $O(10/\log n) = O(1/\log n)$.

The ten regions are sectors or parts of sectors of the circle of radius $\frac{1}{2}x_P$ centered at P. *(Fig. 4.2.)* The central angle of each sector is $\pi/8$. Eight of the regions fall within the sector defined by $\angle CPD$. They must completely cover this sector, but they may otherwise be fixed arbitrarily. To these are added the sector whose upper boundary is PC and the sector whose lower boundary is PD. Since $\angle P'PM$ and $\angle P'PN$ measure $\pi/4$, and $|PM|$, $|PQ|$ and $|PN|$ exceed $\frac{1}{2}x_P$, each of $\triangle PQM$ and $\triangle PQN$ completely contains at least one of these sectors, no matter where P' lies in $\square CDEF$.

The area of the sectors is $\frac{1}{2}(\frac{1}{2}x_P)^2(\pi/8) = (\pi/64)x_P^2 \geq \pi \log\log n/n \geq 2\log\log n/n$. But part of the sector below PC may lie outside $\square ABCD$. In this case let J be the point on PC at distance $\frac{1}{2}x_P$ from P, and let K be the point of intersection of the x-axis and the other side of the sector. Surely the intersection of $\square ABCD$ and the sector contains $\triangle PJK$. But $\text{Area}(\triangle PJK) = \frac{1}{2} \cdot |PJ| \cdot |PK| \cdot \sin(\pi/8) \geq \frac{1}{2} \cdot \frac{1}{2}x_P \cdot y_P \cdot \frac{1}{3} \geq \frac{1}{12}x_P y_P \geq \frac{1}{12}(24\log\log n/n) \geq$


Fig. 4.2. Bounding a truncated sector.

$2\log\log n/n$. Thus we have showed that the intersection of $\square ABCD$ and the sector below PC is itself large enough. The proof is completed by a symmetric argument applied to the sector above PD. $\square$

Theorem 4.6. *Step 2 of Algorithm A requires $O(n\log\log n)$ expected time.*

Proof. The merge region for each merge in Step 2 has height $h = \sqrt{\log n/n}$, thus the area of each

merge region is $O((\sqrt{\log n/n} \cdot \sqrt{\log\log n/n}) + \log n \log\log n/n) = O(\log n \log\log n/n)$. For a particular value of k there are $2^{-k}n/\log n$ disjoint merge regions with total area $O(2^{-k}\log\log n)$. The expected number of sites in these regions is $O(2^{-k}n\log\log n)$. In addition it is expected that at most another $O(n/\log n)$ sites from outside the merge regions will receive new edges. Since by Theorem 2.7 the running time is bounded by a linear function of the number of edges created, and by Corollary 2.2 this is in turn linear in the number of sites receiving new edges, the expected running time of Step 2 is bounded by

$$\sum_{k=0}^{\lg\left(\sqrt{n/\log n}\right)} \left(O(2^{-k}n\log\log n) + O(n/\log n)\right)$$

$$= O(n\log\log n) + O(n) = O(n\log\log n). \ \square$$

Theorem 4.7. *Step 3 of Algorithm A requires $O(n)$ expected time.*

Proof. The merge region for each merge in Step 3 has height $h = 1$ and area $O(\sqrt{\log\log n/n})$. There are $2^{-k}\sqrt{n/\log n}$ merge regions for each k value for a total area of $O(2^{-k}\sqrt{\log\log n/\log n}) = O(2^{-k})$ containing $O(2^{-k}n)$ sites in the expected case. Reasoning as for Step 2, we find the running time in the expected case to be bounded by

$$\sum_{k=0}^{\lg\left(\sqrt{n/\log n}\right)} \left(O(2^{-k}n) + O(n/\log n)\right) = O(n). \ \square$$

Adding together the expected time for each step we have:

Theorem 4.8. *Algorithm A constructs $DT(S)$ in $O(n\log\log n)$ expected time.*

5. Extension to the L_p Metrics.

In this section we show that the algorithm of Section 3 and its analysis extend to the L_p metrics for $1 < p \leq \infty$. We are able to make only a small improvement for the L_1 case.

The L_p metric is defined by the distance function $d_p(P,Q) = (|x_P - x_Q|^p + |y_P - y_Q|^p)^{1/p}$ for $1 \leq p < \infty$ and $d_\infty(P,Q) = \max(|x_P - x_Q|, |y_P - y_Q|)$ for $p = \infty$. The L_2 metric is the

usual Euclidean metric. The L_1 metric is some-
times called the *rectilinear* or *Manhattan metric*.
Hwang [H], Lee & Wong [LW], and Lee [L] present
$O(n \log n)$ divide-and-conquer algorithms for con-
structing the Voronoi diagrams in the L_1, L_1
and L_∞, and general L_p metrics respectively. In
these metrics each Voronoi polygon is star-shaped
with a nucleus at the associated site, but it is
not necessarily convex. The straight-line dual of
the Voronoi diagram is called the Delaunay tri-
angulation and denoted $DT_p(S)$, although in fact
$DT_1(S)$ and $DT_\infty(S)$ do not generally triangu-
late the convex hull of S.

Lee [L] shows that Lemma 2.3, Corollaries
2.4 and 2.5, and (essentially) Theorem 2.6 hold
for $1 < p < \infty$ when the word "circle" is taken to
mean the locus of points at a constant distance in
the L_p metric from some given point. From these
Theorems 2.7 and 3.1 follow. He also proves the
following useful lemma.

*Lemma 5.1. The locus of points equidistant from
two given points, called their bisector, is either
nonincreasing or nondecreasing.*

We can now extend the average-case analysis
to $1 < p < \infty$.

*Theorem 5.2. Algorithm A constructs $DT_p(S)$
in $O(n \log \log n)$ expected time for $1 < p < \infty$.*

Proof. The analysis of Section 4 lacks only an L_p
version of Lemma 4.3, which we now provide:

*Lemma 5.3. Let P, Q, and P' be three points
lying along a line in that order. Let M and N be
the other vertices of the square with diagonal PQ.
Then for $1 \leq p \leq \infty$ any L_p circle passing through
P and P' completely encloses either $\triangle PQM$ or
$\triangle PQN$ or both.*

Proof. Let M' and N' be the other vertices of the
square with diagonal PP' lying on the same side
of PP' as M and N respectively. Since $\triangle PP'M'$
and $\triangle PP'N'$ contain $\triangle PQM$ and $\triangle PQN$ re-
spectively, it suffices to show that the circle con-
tains either M' or N'. *(Fig. 5.1.)* Without loss
of generality we consider only the case in which
$a \geq b \geq 0$, $P = (0, 2b)$, $P' = (2a, 0)$, and
$M' = (a + b, a + b)$, and show that the center
of the circle C must lie nearer M' than P if it lies
above the line through P and P'.

The triangle inequality implies that the bi-

Fig. 5.1. Bounding the area of L_p circles.

sector of P and P' passes no nearer to P than
$(a^p + b^p)^{1/p}$ at $R = (a, b)$. If $C = R$, M' clearly
lies on the perimeter of the circle. Since both
M' and R lie on the bisector, it is nondecreasing
rather than nonincreasing. Thus $a \leq x_C \leq a + b$
and $b \leq y_C \leq a + b$ if C lies on the bisector be-
tween R and M', and $d_p(C, M') = (|a+b-x_C|^p +
|a + b - y_C|^p)^{1/p} \leq (b^p + a^p)^{1/p} = d_p(R, P) \leq
d_p(C, P)$. Beyond M', C must satisfy $x_C \geq a + b$
and $y_C \geq a + b$, so $d_p(C, M') = ((x_C - (a + b))^p +
(y_C - (a + b))^p)^{1/p} \leq (x_C^p + (y_C - 2b)^p)^{1/p} =
d_p(C, P)$. $\square$

It is the L_1 and L_∞ Voronoi diagrams which
are of most interest in applications. Unfortu-
nately, the implementation and analysis of these
cases is complicated by the fact that some con-
vex hull edges may be omitted from the Delaunay
triangulation. For these cases, the Guibas-Stolfi
merge procedure must be modified along the lines
of Lee's dual algorithm to search downward for
the endpoints of the lower convex hull edge then
back upward for the endpoints of the lowest De-
launay edge crossing the dividing line. Thus some
sites examined in the search do not receive new
edges during the merge, and Theorem 2.7, which
relates the time spent on these lower-boundary-
edge searches to the total number of edges cre-
ated, no longer holds. The analysis of Section 4
still bounds the expected number of edges cre-
ated, but not the overall running time.

We call the edges of the infinite face of the
Delauanay triangulation *boundary edges* and their
endpoints *boundary sites*. We show that the L_∞
boundary sites of a set of sites contained in an
orthogonal rectangle in the unit square all proba-
bly lie near the boundary of the rectangle. There
are therefore so few of them in the expected case
that searching through them does not dominate
the cost of the merge step.

Lemma 5.4. *A site $P = (x_P, y_P)$ is a boundary site in $DT_\infty(S)$ only if at least one of the quarter-planes defined by the lines $x = x_P$ and $y = y_P$ is site-free.*

Proof. The result follows immediately from Lee's observation that an edge is a boundary edge in $DT_\infty(S)$ if and only if one of the two quarter-planes defined by horizontal and vertical rays passing though the endpoints is site-free. $\square$

Theorem 5.5. *Let $\square ABCD$ be an orthogonal rectangle lying inside $\mathcal{U}$. Then there is a boundary region $\mathcal{B} \subset \square ABCD$ such that*
a) $\mathrm{Area}(\mathcal{B}) = O(\log n \log \log n / n)$
b) if $P \in (S \cap \square ABCD - \mathcal{B})$ then $p = \Pr\{P$ is a boundary site in $DT_\infty(S)\} = O(1/\log n)$.

Proof. For convenience, we use a coordinate system with origin at C, B on the positive x-axis, and D on the positive y-axis. We now verify that conditions a) and b) are satisfied by

$$\mathcal{B} = \square ABCD \cap \big\{ (x, y) \mid (xy \leq 2 \log \log n / n) \\ \vee (x(h - y) \leq 2 \log \log n / n) \\ \vee ((w - x)y \leq 2 \log \log n / n) \\ \vee ((w - x)(h - y) \leq 2 \log \log n / n) \big\}$$

The bound on $\mathrm{Area}(\mathcal{B})$ is straightforward since $h, w \leq 1$.

Without loss of generality suppose that $P = (x_P, y_P)$ lies in the lower left quarter of $\square ABCD$ but outside $\mathcal{B}$. Then the smallest of the four rectangles into which $\square ABCD$ is partitioned by the vertical and horizontal lines through P has area $x_P y_P \geq 2 \log \log n / n$. By Lemma 4.4 the probability of its being site-free is at most $O(1/\log n)$. Therefore by Lemma 5.4 the probability that P is on the boundary is at most $O(4/\log n) = O(1/\log n)$. $\square$

We can now bound the running time of Algorithm A for the L_∞ metric.

Theorem 5.6. *Algorithm A constructs $DT_\infty(S)$ in $O(n \log \log n)$ expected time.*

Proof. The analysis of Step 1 in Theorem 4.2 remains valid, since the worse-case performance of the L_1 / L_∞ merge procedure is $O(n \log n)$. Since the boundary region of any rectangle is smaller than its merge region, the calculations of Theorems 4.6 and 4.7 for Steps 2 and 3 apply to the number of sites examined in boundary-edge searches as well as to the number of edges cre-

ated. Therefore the overall bound of Theorem 4.8 applies as well. $\square$

The L_1 case is more difficult. It follows from Lemma 5.4 and the well-known correspondence between the L_1 and L_∞ metrics that a site P is an L_1 boundary site if and only if one of the quarter-planes defined by the lines through P with slopes $+1$ and -1 is site-free. The L_1 analog of Theorem 5.5 states that all points within a (Euclidean) distance of $\sqrt{2 \log \log n / n}$ of the boundary fall within the boundary region. Analyzing the number of sites examined in lower-boundary-edge searches in Steps 2 and 3 yields an $O(n\sqrt{\log n \log \log n})$ bound, which is only a marginal improvement to the original $\Theta(n \log n)$ expected running time.

6. Experimental Results.

A variation of Algorithm A for the Euclidean metric has been implemented for practical evaluation. In this variation, the set of n sites is divided into $\sqrt{n / \log n}$ equal subsets by lines perpendicular to the y-axis. Then the Guibas-Stolfi algorithm, which divides with lines perpendicular to the x-axis, is applied to each subset. Finally the results are merged in pairs as in Step 3 of Algorithm A. This variation is somewhat easier to implement but more difficult to analyze. Intuitively, we expect its behavior to differ little from that of Algorithm A for reasonably large n.


Fig. 6.1. The algorithms compared.

This variation and the original Guibas-Stolfi algorithm were run on inputs generated by drawing sites from the uniform distribution in the unit square. Twenty inputs of size 2^k were generated for $4 \leq k \leq 16$. The results are summarized in Fig. 6.1, which plots the mean number of edges created per site as a function of n, the number

of sites. The small variance can be safely ignored. Our measurements for the original algorithm match closely those of Ohya, Iri & Murota. It is clear that the modified algorithm is significantly faster for all but the smallest values of n.


Fig. 6.2. Estimating the constant factor.

Fig. 6.2 is useful in estimating the constant factor c in the upper bound $cn \log \log n$ on expected running time. If c is asymptotically less than 1.77 as Fig. 6.2 suggests, the number of edges created by the algorithm would be less than $6n$ for $n \leq \exp(\exp(6.0/1.77)) \approx 7 \times 10^{12}$. Since about $3n$ edges are required in the final diagram, we conjecture that the running time is no more than twice optimal for n in this range.

7. Conclusions.

We have showed how the average running time of the Guibas-Stolfi algorithm for constructing the Delaunay triangulation can be improved dramatically to $O(n \log \log n)$ for a large class of distributions of the sites. We believe the improved algorithm is competitive with linear-time incremental algorithms even for relatively large problems, but cannot state this unequivocably without further theoretical and experimental study of the algorithm's performance on non-uniform distributions and direct experimental comparison with the incremental algorithms.

We have also showed that the improvement extends to the L_p version of the algorithm for $1 < p \leq \infty$. It would naturally be desirable to extend it to the L_1 case.

8. Acknowledgements.

Thanks to Jon Webb, whose code for the Guibas-Stolfi algorithm was the starting-point

9. References.

[BWY] J. L. Bentley, B. W. Weide & A. C. Yao, "Optimal expected-time algorithms for closest point problems," *ACM Trans. Math. Software* **6** (1980), 563–580.

[F] S. Fortune, "A sweepline algorithm for Voronoi diagrams," *these proceedings*.

[GS] L. J. Guibas & J. Stolfi, "Primitives for the manipulation of general subdivisions and the computation of Voronoi diagrams," *ACM Trans. Graphics* **4** (1985), 74–123.

[H] F. K. Hwang, "An $O(n \log n)$ algorithm for rectilinear minimal spanning trees," *J. ACM* **26** (1979), 177–182.

[L] D. T. Lee, "Two-dimensional Voronoi diagrams in the L_p-metric," *J. ACM* **27** (1980), 604–618.

[LS] D. T. Lee & B. Schachter, "Two algorithms for constructing Delaunay triangulations," *Int. J. Comput. Inform. Sci.* **9** (1980), 219–242.

[LW] D. T. Lee & C. K. Wong, "Voronoi diagrams in L_1 (L_∞) metrics with two-dimensional storage applications," *SIAM J. Comput.* **9** (1980), 200–211.

[M] A. Maus, "Delaunay triangulation and the convex hull of n points in expected linear time," *BIT* **24** (1984), 151–163.

[OIM] T. Ohya, M. Iri & K. Murota, "Improvements of the incremental methods for the Voronoi diagram with computational comparison of various algorithms," *J. Operations Research Soc. Japan* **27** (1984), 306–336.

[S] R. Sibson, "Locally equiangular triangulations," *Comput. J.* **21** (1978), 243-245.

[SG] R. Sibson & P. J. Green, "Computing Dirichlet tessellations in the plane," *Comput. J.* **21** (1978), 168–173.

[Sh] M. I. Shamos, *Computational Geometry,* Ph.D. Thesis (1978), Yale University.